

Empathica AI: Brand-Aligned Response Generator

Amazon Customer Support | Generative AI Project

A Project Report

Submitted by

Venkata Kartheek Sanka (AP23110010750)

Sumanth Oruganti (AP23110010786)

Guna Vardhan Byraju (AP23110010744)

SEM:VI

SECTION: K

BACHELOR OF TECHNOLOGY

in

COMPUTER SCIENCE AND ENGINEERING

of

SCHOOL OF ENGINEERING AND SCIENCES



SRM
UNIVERSITY AP

—Andhra Pradesh

SRM University-AP, Neerukonda, Andhra Pradesh 522240

May 2026

TABLE OF CONTENTS

Chapter 1: Project Overview	3-5
1.1 Problem Statement	3
1.2 Solution	4
1.3 Real World Example	4
1.4 Development Achievements	5
Chapter 2: System Architecture	6-7
2.1 Dual Model Pipeline	6
2.2 Full Stack Design	7
2.3 Technology Stack	7
Chapter 3: Dataset & Data Preparation	8-9
3.1 Dataset Selection	8
3.2 Data Preparation Pipeline	8-9
3.3 Dataset Statistics	9
Chapter 4: Model Architecture	10-12
4.1 Base Model — Flan-T5	10
4.2 Fine-Tuning Using LoRA	11
4.3 Emotion Detection Model	12
Chapter 5: Training Process	13-14
5.1 Tokenization	13
5.2 Training Configuration	13-14
5.3 Training Results	14
Chapter 6: Inference Pipeline	15-16
6.1 Complete Inference Flow	15
6.2 Brand Tone Examples	16
Chapter 7: Evaluation & Results	17-18
7.1 ROUGE Evaluation	17
7.2 Before vs After Comparison	18
Chapter 8: Screenshots of the Application	19
Chapter 9: Backend Implementation	20-21
9.1 FastAPI Application	20
9.2 API Endpoints	21
Chapter 10: Frontend Implementation	22
Chapter 11: How to Run the Project	23
Chapter 12: Key Concepts Glossary	24
Chapter 13: Conclusion & Future Work	24
REFERENCES	24

Chapter 1: Project Overview

1.1 Problem Statement

Daily, e-commerce giants like Amazon have millions of consumers making complaints, such as delays in receiving packages, the wrong package arriving, billing errors or refund issues. The traditional methods of supplying human written responses to these consumers has three primary limitations:

- **Scale:** There is no way on the planet to have enough human agents that can respond to 1 million complaints a day consistently in terms of quality of response.
- **Inconsistency:** Because there are different support agents responding to complaints, you have various tonal qualities in the agent's response, which disrupts the integrity of the brand.
- **Robotic Automation:** Basic chatbot responses are cold, templated and are void of any emotional content, which typically makes consumers even angrier with their company.

The Core Problem In One Sentence: “How do we create millions of empathetic (i.e., emotional), brand consistent customer service responses via an automated system, at scale?”

1.2 Solution

Empathica AI Empathica AI is a cutting-edge Generative AI system designed with a two-model pipeline:

- **Step 1 — Emotion Detection:** The DistilRoBERTa Classifier is an easily deployable Emotion Detector that accepts the customer's complaint and determines the emotion behind it (I.e., angry, sad, scared, etc.).
- **Step 2 — Emotion-Based Response Generation:** The Flan-T5 Model is a Fine-tuned Generative Model trained with LoRA on a dataset of 3000 support conversations in an Amazon format. The generation of an empathic response is conditioned by the detected emotion and the desired tone of the brand.

1.3 Real World Example

A customer complains about their order not having arrived after two weeks of being placed and that they have not received any assistance (Sad emotion detected). The Empathica AI model generated the following response: "Thank you [the customer] for reaching out to Amazon Support. I understand your frustration with the fact that your order has not yet arrived, and please know I'm here to help you find a resolution. If you have any other questions, I am happy to help you." (Friendly Tone).

1.4 Development Achievements


Metric	Our Score	Good Threshold	Status
ROUGE-1	0.4495	> 0.30	Excellent (+49%)
ROUGE-2	0.2037	> 0.15	Good (+35%)
ROUGE-L	0.2991	> 0.25	Good (+19%)
Trainable Params	0.71%	< 5%	Efficient
Training Time	~8 min	GPU T4	Fast

Chapter 2: System Architecture

Section 2.1 The Dual Model Pipeline of the Hybrid Platform


The system itself consists of two Artificial Intelligence (AI) models that process customer complaints into empathetic responses by going through a four-stage flow. The flow includes the following steps:

Customer Complaint (text input) -> Emotion Detector (DistilRoBERTa) -> Prompt Builder (Tone and/or Emotion) --> Flan-T5 + LoRA (Fine-tuned model) --> Empathetic Response (output).



Section 2.2 Full Stack Design

Layer	Components
Frontend Layer	index.html demo.html compare.html style.css script.js
FastAPI Backend (main.py)	/generate endpoint /health endpoint static file serving
Emotion Detector	j-hartmann/emotion-english-distilroberta-base 7 emotion classes
Fine-Tuned Flan-T5 + LoRA	google/flan-t5-base LoRA rank=16 Trained on 2,700 samples
Dataset	Bitext Customer Support Dataset HuggingFace 3,000 Amazon-style samples



2.3 Technology Stack

Component	Technology	Purpose
Core Language	Python 3.10+	All backend logic and model inference
Base LLM	google/flan-t5-base	250M parameter seq2seq model
Fine-Tuning Method	PEFT + LoRA	Parameter-efficient adaptation
Emotion Detection	DistilRoBERTa	7-class emotion classification
ML Framework	HuggingFace Transformers	Model loading and training
Training Library	Seq2SeqTrainer	Fine-tuning pipeline
Dataset	Bitext (HuggingFace)	3,000 Amazon-style support pairs
Backend API	FastAPI + Uvicorn	REST API serving
Frontend	HTML5 + CSS3 + JS	Web user interface
Evaluation	ROUGE Score	Response quality measurement

Chapter 3: Dataset & Data Preparation

3.1 Dataset Selection

We used the Bitext Customer Support LLM Chatbot Training Dataset available on HuggingFace:

Dataset Property	Details
Name	Bitext-customer-support-llm-chatbot-training-dataset
Source	HuggingFace Datasets
Total Tokens	3.57 Million
Format	instruction (complaint) + response (reply) pairs
License	CDLA-Sharing 1.0 (Free for research)

3.1.1 What is the purpose of this dataset?

- Structured in pair formats — fits extremely well for sequence-to-sequence approaches
- Includes complete coverage of the defined categories of complaints related to Amazon (including, but not limited to: order questions, refund questions, delivery questions, payment questions)
- Created by computational linguistic experts with high standards; no data that is irrelevant or useless
- Can be loaded directly via one line of Python code through HuggingFace datasets.

Step 1 — Load Dataset

```
from datasets import load_dataset

dataset = load_dataset(
    'bitext/Bitext-customer-support-llm-chatbot-training-dataset'
)
df = pd.DataFrame(dataset['train'])
```

Step 2 — Filter Amazon-Relevant Categories

```
amazon_categories = [
    'ORDER', 'REFUND', 'DELIVERY', 'INVOICE',
    'PAYMENT', 'ACCOUNT', 'CANCELLATION_AND_RETURN'
]
df_amazon = df[df['category'].isin(amazon_categories)].copy()
```

Step 3 — Clean and Brand-Align

```
def clean_text(text):
    text = str(text).strip()
    text = text.replace('\n', ' ')
    text = ' '.join(text.split())
    return text
```

```
def add_amazon_style(response):
    return f'Thank you for contacting Amazon Support. {response}'

df_amazon['instruction'] = df_amazon['instruction'].apply(clean_text)
df_amazon['response'] = df_amazon['response'].apply(clean_text)
df_amazon['response'] = df_amazon['response'].apply(add_amazon_style)
```

Step 4 — Train/Test Split

```
df_sample = df_amazon.sample(n=3000, random_state=42).reset_index(drop=True)
train_size = int(0.9 * len(df_sample))

train_df = df_sample[:train_size] # 2,700 samples
test_df = df_sample[train_size:] # 300 samples
```

3.3 Dataset Statistics

Split	Samples	Percentage
Training Set	2,700	90%
Test Set	300	10%
Total	3,000	100%

Chapter 4: Model Architecture

4.1 Base Model — Google Flan-T5

4.1.1 What is Flan-T5?

Flan-T5 (Fine-tuned Language Net — Text-to-Text Transfer Transformer) is a sequence-based natural language processing model developed by Google that has been pre-trained on large amounts of text from the internet and then fine-tuned with instructions from other NLP tasks.

Property	Value
Model Name	google/flan-t5-base
Architecture	Encoder-Decoder (T5)

Total Params	249,347,328 (~250M)
Task Type	Sequence-to-Sequence (Text Generation)
Input Length	Up to 512 tokens
Pre-training	Instruction-tuned on 1,800+ tasks

4.1.2 Why Flan-T5 over GPT or LLaMA?

Model	Params	GPU Needed	Cost	Choice
GPT-4	~1T	API only	Paid	Not selected
LlaMA-7B	7B	A100 80GB	Expensive	Not selected
Mistral-7B	7B	A100 40GB	Expensive	Not selected
Flan-T5-Base	250M	T4 16GB	Free	Selected


4.2 Fine-Tuning Using LoRA

4.2.1 The Problem with Complete Fine-Tuning

Complete fine-tuning of a model with 250 million parameters necessitates the simultaneous modification of all 250 million weights, which requires very large amounts of GPU memory, a lot of compute power to train, and complete hours of training.

4.2.2 What Is LoRA?

Low-Rank Adaptation (LoRA) involves maintaining the original model parameters and using additional trainable adapter matrices that are included with the existing layers within the model. This provides a means to allow the model to learn domain-specific characteristics while only modifying 1% or less of the original model parameters.



4.2.3 LoRA Configuration

```

from peft import LoraConfig, TaskType

lora_config = LoraConfig(

```

```
task_type = TaskType.SEQ_2_SEQ_LM,
r = 16, # Rank: number of trainable dimensions
lora_alpha = 32, # Scaling factor (2x of r is standard)
lora_dropout = 0.1, # 10% dropout prevents overfitting
target_modules = ['q', 'v'] # Apply to attention Q and V matrices
)
```

4.2.4 Parameter Efficiency Results

Parameter Type	Count	Percentage
Total Parameters	249,347,328	100.00%
Frozen Parameters	247,577,856	99.29%
Trainable (LoRA)	1,769,472	0.71%

Key Insight: By training only 0.71% of the model's parameters, we achieve Amazon-style empathetic response generation. This is the power of LoRA — maximum result with minimum compute!

4.3 Emotion Detection Model

4.3.1 Model Details

Property	Value
Model	j-hartmann/emotion-english-distilroberta-base
Base Architecture	DistilRoBERTa (66M params)
Task	Multi-class text classification
Output Classes	7 emotion labels
Input Limit	512 tokens

4.3.2 Emotion Classes

Label	Emotion	Example Trigger
[ANGER]	Anger	"Fix this immediately!"
[SAD]	Sadness	"Nobody is helping me"
[FEAR]	Fear	"I'm scared I lost my money"
[JOY]	Joy	"Thank you so much!"

[SURP]	Surprise	"I can't believe this happened"
[DISG]	Disgust	"This is completely unacceptable"
[NEUT]	Neutral	"I need to cancel my order"

Chapter 5: Training Process

5.1 Tokenization

Before training, all text is converted to numbers (tokens) that the model understands. Maximum input and target lengths are both set to 128 tokens.

```
MAX_INPUT_LENGTH = 128
MAX_TARGET_LENGTH = 128

def preprocess_function(examples):
    inputs = [
        f'Generate empathetic Amazon support response: {inst}'
        for inst in examples['instruction']
    ]
    model_inputs = tokenizer(
        inputs, max_length=MAX_INPUT_LENGTH,
        truncation=True, padding='max_length'
    )
    labels = tokenizer(
        text_target=examples['response'],
        max_length=MAX_TARGET_LENGTH,
        truncation=True, padding='max_length'
    )
    label_ids = []
    for label in labels['input_ids']:
        label_ids.append([
            -100 if token == tokenizer.pad_token_id else token
            for token in label
        ])
    model_inputs['labels'] = label_ids
    return model_inputs
```

5.2 Training Configuration

```
training_args = Seq2SeqTrainingArguments(
    output_dir = './empathica-amazon-model-v3',
    num_train_epochs = 3,
    per_device_train_batch_size = 4,
    per_device_eval_batch_size = 4,
    warmup_steps = 100,
    weight_decay = 0.01,
    logging_steps = 25,
    eval_strategy = 'epoch',
    save_strategy = 'epoch',
    load_best_model_at_end = True,
    predict_with_generate = True,
```


```
fp16          = False, # caused nan loss on T4
report_to     = 'none'
)
```

5.2.1 Hyperparameter Rationale

Hyperparameter	Value	Reason
Epochs	3	Enough to converge without overfitting on 2,700 samples
Batch Size	4	Fits within T4 GPU 16 GB VRAM
Warmup Steps	100	Prevents large gradient updates at the start
Weight Decay	0.01	L2 regularisation to prevent memorising training data
fp16	False	Mixed precision caused NaN loss on T4; disabled

5.3 Training Results

Epoch	Training Loss	Validation Loss	Trend
1	1.9168	1.6898	Improving
2	1.8328	1.5584	Improving
3	1.7736	1.5264	Converging



Observation: Both losses decrease consistently across all 3 epochs, confirming the model is learning without overfitting.

Chapter 6: Inference Pipeline

6.1 Complete Inference Flow

```
def generate_response(complaint: str, brand_tone: str = 'friendly'):

    # STEP 1: Detect emotion in complaint
    emotion = detect_emotion(complaint)

    # STEP 2: Select tone-specific prompt
    tone_prompts = {
        'friendly' : 'Generate a friendly and empathetic Amazon support response:',
        'formal'   : 'Generate a formal and professional Amazon support response:',
        'apologetic' : 'Generate a deeply apologetic Amazon support response:'
    }

    # STEP 3: Build the full prompt
    prompt = (
        f"{tone_prompts[brand_tone]} "
        f"Customer complaint: {complaint} "
        f"Detected emotion: {emotion}"
    )

    # STEP 4: Tokenize
    inputs = tokenizer(prompt, return_tensors='pt',
                       max_length=128, truncation=True).to(model.device)

    # STEP 5: Generate response using beam search
    with torch.no_grad():
        outputs = model.generate(
            **inputs,
            max_new_tokens    = 150,
            num_beams         = 4,
            early_stopping    = True,
            no_repeat_ngram_size = 2
        )

    # STEP 6: Decode tokens back to text
    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    return emotion, response
```

6.2 Brand Tone Examples

Tone	Sample Output for "My order is 2 weeks late"
Friendly	"Thank you for contacting Amazon Support! We completely understand how frustrating this must be, and we're so sorry for the delay. We're looking into this right away for you!"
Formal	"Thank you for contacting Amazon Support. We acknowledge your inquiry regarding the delay in delivery. Our logistics team has been notified and will provide an update within 24 business hours."

Apologetic	"Thank you for contacting Amazon Support. We sincerely and deeply apologize for this unacceptable delay. This is entirely our responsibility and we are taking immediate action to resolve your issue."
------------	---

Chapter 7: Evaluation & Results

7.1 Evaluation of the ROUGE Score

7.1.1 What is ROUGE? ROUGE (Recall-Oriented Understudy for Gisting Evaluation)

checks how similar a generated text is to a reference (ideal) text. It is the standard way to measure how well a text generation task is done.

Metric	What it Measures	Formula
ROUGE-1	Matching individual words (unigrams)	common words / reference words
ROUGE-2	Matching word pairs (bigrams)	common bigrams / reference bigrams
ROUGE-L	Longest common subsequence	LCS length / reference length

7.1.2 Our Results

ROUGE Metric	Our Score	Threshold	Result
ROUGE-1	0.4495	> 0.30	Excellent (+49%)
ROUGE-2	0.2037	> 0.15	Good (+35%)
ROUGE-L	0.2991	> 0.25	Good (+19%)

7.2 Before vs After Fine-Tuning Comparison



Complaint	Base Flan-T5 (No Fine-Tuning)	Empathica AI (Fine-Tuned)
"My order hasn't arrived in 2 weeks!"	"My order hasn't arrived in 2 weeks." (just repeats!)	"Thank you for contacting Amazon Support. I understand your frustration. I'm here to assist you immediately."
"Charged twice for same item"	"I was charged twice for the same item."	"Thank you for contacting Amazon Support. We understand your concern and will process the refund immediately."
"Cancel button not working"	"Is there anything else I can help you with?"	"Thank you for contacting Amazon Support. I will guide you through the cancellation process right now."


Important Point: The base model either repeats the complaint or gives an answer that has nothing to do with it. Empathica AI always gives warm, understanding, and brand-aligned answers, which shows how useful LoRA fine-tuning is.

Chapter 8: Screenshots of the Application

Run your project on your own computer (uvicorn backend.main:app --reload --port 8000), take the screenshots below, and put them in the boxes below.


8.1 Screenshot 1: The Landing Page (System Architecture)






Why this screenshot is important: Shows evaluators that you really understand the technical architecture in a way that is professional and ready for consumers.

8.2 Screenshot 2 — Live Inference Engine (Action Shot)



Why this screenshot is important: It shows the main product. It shows the workspace, the detected emotion badge, and the output from the fine-tuned model that is of high quality.

8.3 Screenshot 3 — Model Comparison (Before & After)



Model Comparison

Watch our fine-tuned Empathica AI model run head-to-head against the raw base FLAN-T5 model to see the dramatic difference in empathy and tone alignment.

CUSTOMER COMPLAINT

My order is 3 weeks late and Amazon customer service keeps telling me to wait .This is unacceptable

99 / 512

TARGET TONE

Friendly Formal **Apologetic**

QUICK EXAMPLES

Order 3 weeks late, no support Charged twice, need refund

COMPARE BOTH MODELS

VS

BASE FLAN-T5 No Fine-Tuning

Amazon has apologised for the customer complaint.

EMPATHICA AI LoRA Fine-Tuned

Thank you for contacting Amazon Support. I understand your frustration that your order is 3 weeks late and you're told to wait. We understand that this is unacceptable, and we are committed to providing you with the best service possible. To ensure this, I'm here to assist you in resolving any issues you may have. Let me know if you have any further questions or concerns.

Anger

This screenshot is important because it is the most important technical proof in your report. It shows in a clear way that your fine-tuning worked.

8.4 Screenshot 4 — REST API Documentation

REST API Reference

Integrate Empathica AI's fine-tuned empathy models directly into your platform via our local FastAPI backend. Seamless JSON responses, automatic validation, and sub-second inference.

CONTENTS

Overview

GET /health

POST /generate

Code Examples

Error Handling

Overview

Empathica AI provides a simple REST API built with FastAPI. The API runs locally on port 8000.

BASE URL: `http://localhost:8000`

`GET` /health

Returns system status and the device the models are running on. Use this to confirm the server started and models loaded correctly.

Response

```
JSON Response - 200 OK
{
  "status": "running",
  "device": "cpu",
  "model": "Empathica AI v1.0"
}
```

Why this screenshot is important: Shows that the project is more than just a UI demo; it also has a backend service that works and is ready to be integrated.

Chapter 9: Backend Implementation

9.1 FastAPI Application

```
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
from fastapi.staticfiles import StaticFiles
from fastapi.responses import FileResponse
from pydantic import BaseModel

app = FastAPI(title='Empathica AI')

app.add_middleware(
    CORSMiddleware,
    allow_origins=['*'], allow_methods=['*'], allow_headers=['*']
)

app.mount('/static', StaticFiles(directory='frontend'), name='static')

class ComplaintRequest(BaseModel):
    complaint : str
    tone      : str = 'friendly'

@app.get('/')
def home():
    return FileResponse('frontend/index.html')

@app.post('/generate')
def generate(req: ComplaintRequest):
    emotion, response = generate_response(req.complaint, req.tone)
    return {'emotion': emotion, 'response': response, 'tone': req.tone}

@app.get('/health')
def health():
    return {'status': 'running', 'device': device}
```

9.2 API Endpoints

Method	Endpoint	Type	Description
GET	/	HTML	Serves the landing page
POST	/generate	JSON API	Main inference endpoint
GET	/health	JSON	Server health check
GET	/demo	HTML	Live demo workspace
GET	/compare	HTML	Before/after comparison
GET	/api-docs	HTML	API documentation page

9.2.1 API Request / Response Schema

Request (POST /generate):

```
{
  "complaint" : "My order hasn't arrived in 2 weeks!",
  "tone"      : "friendly" // Options: friendly | formal | apologetic
}
```

Response:

```
{
  "emotion" : "sadness",
  "response" : "Thank you for contacting Amazon Support. I understand
               your frustration that your order hasn't arrived...",
  "tone"     : "friendly"
}
```

Chapter 10: Frontend Implementation

10.1 Application Pages

File	Route	Purpose
index.html	/	Landing page with architecture overview
demo.html	/demo	Live inference workspace
compare.html	/compare	Base model vs fine-tuned comparison
docs.html	/api-docs	REST API reference documentation
style.css	static	Shared design system
script.js	static	All interaction logic

10.2 Important UI Components

- Tone Selection Options — three button options (Friendly/Formal/Apologetic) highlight chosen state.
- Emotion Badge — always display an emotion using colour, i.e. Anger/Sadness/Neutral.
- Example Complaints — allows the user to quickly fill out a sample of various types of complaints.
- Copy Button — simply click a button to copy the generated response to clipboard.
- Statistics Bar Footer — key statistics displayed at the footer of the Application, e.g. 3000+ total samples, 0.44 IMEM, 7 total Emotions detected and 3 Tones detected.
- Split Screen Comparison Tool — users can see and compare the responses generated by the different models on the /compare Page side by side by side, and compare their individual results.

Chapter 11: How to Run the Project

11.1 Project Structure

```
empathica-ai/
|
|-- model/
|   |-- empathica-amazon-final/  <- Fine-tuned model files
|   |-- adapter_config.json
|   |-- adapter_model.bin
|   |-- tokenizer files...
|
|-- backend/
|   |-- main.py                  <- FastAPI application
|
|-- frontend/
|   |-- index.html               <- Landing page
|   |-- demo.html               <- Live demo
|   |-- compare.html            <- Comparison page
|   |-- docs.html               <- API documentation
|   |-- style.css               <- Design system
|   |-- script.js               <- Frontend logic
|
|-- requirements.txt
|-- START_EMPATHICA.bat         <- Windows launcher
```

11.2 Installation & Setup

```
# Step 1: Enter the project directory
cd empathica-ai

# Step 2: Install Python dependencies
pip install -r requirements.txt

# Step 3: Run the FastAPI server
uvicorn backend.main:app --reload --port 8000

# Step 4: Open in browser
# http://localhost:8000
```

11.3 Requirements

```
transformers>=4.40.0
torch>=2.0.0
peft>=0.10.0
accelerate>=0.27.0
datasets>=2.18.0
fastapi>=0.110.0
uvicorn>=0.27.0
sentencepiece>=0.2.0
evaluate>=0.4.0
```

rouge-score>=0.1.2

Chapter 12: Key Concepts Glossary

Term	Full Form	Simple Explanation
LLM	Large Language Model	AI model trained on massive text data to understand and generate language
PEFT	Parameter Efficient Fine-Tuning	Fine-tune only a tiny portion of the model's weights
LoRA	Low-Rank Adaptation	Add small trainable matrices alongside frozen model weights
Seq2Seq	Sequence to Sequence	Input text → Output text generation architecture
Tokenization	—	Converting words and characters to numbers for AI processing
Epoch	—	One complete pass through the entire training dataset
Loss	—	Measure of how wrong the model is; lower is better
ROUGE	Recall-Oriented Understudy for Gisting Evaluation	Standard metric to measure text generation quality
Beam Search	—	Smart decoding strategy exploring multiple output candidates
Fine-tuning	—	Teaching a pre-trained model new domain-specific skills
Inference	—	Using a trained model to generate predictions on new inputs
Embedding	—	Numerical vector representation of text meaning
Attention	—	Mechanism allowing the model to focus on relevant input words

13.1 Summary

Through this chapter we have investigated how Empathica AI demonstrates a method of using a parameter-efficient fine-tuned language model to be able to scale up generating empathetic and brand-consistent customer support responses. The combination of employing:

- LoRA fine-tuning done on Flan-T5 where only 0.71% of the total parameters of the model were tuned
- Using the DistilRoBERTa for Emotion detection (providing context for generating the appropriate response)
- Using a defined Brand Tone e.g. (Friendly / Formal / Apologetic)
- Utilizing FastAPI and polished Web UI for production deployment

With these implementations we achieved a ROUGE-1 score of 0.4495-49% above the threshold defined in the industry (0.30) from a testing set of 300 actual customer support conversations.13.2 What Makes This Project Stand Out

Feature	Basic Project	Empathica AI
Base LLM used	Yes	Yes
Fine-tuned with LoRA	No	Yes
Emotion detection	No	Yes
Brand tone control	No	Yes
ROUGE evaluation	No	Yes
Before/After proof	No	Yes
Production API	No	Yes
Complete Web UI	No	Yes

13.3 Future Improvements

- Multilingual Support: Expansion to Hindi and Indian languages by employing mT5 or IndicBERT
- Real-Time Streaming: Token streaming implementation for instant response time in the GUI
- Semantic Evaluation via BERTScore: Perform evaluation of response quality along with ROUGE
- Feedback-Based Tuning: Gather thumbs up/down responses from the users for continuous fine-tuning
- Larger Model: Upgrade to Flan-T5-Large or Mistral-7B models for better responses
- Dockerized Applicability: Containerize application for effortless deployment in the cloud using AWS/GCP

References

1. E. J. Hu, Y. Shen, P. Wallis, et al.,
"LoRA: Low-Rank Adaptation of Large Language Models," arXiv:2106.09685, 2021.
2. H. W. Chung, L. Hou, S. Longpre, et al.,
"Scaling Instruction-Finetuned Language Models," arXiv:2210:
<https://huggingface.co/datasets/bitext>
4. J. Hartmann,
"Emotion English DistilRoBERTa-base," HuggingFace Models, 2022.
Available: <https://huggingface.co/j-hartmann/emotion-english-distilroberta-base>
5. T. Wolf, L. Debut, V. Sanh, et al.,
"Transformers: State-of-the-Art Natural Language Processing," EMNLP, 2020.